

Yeqiao Fu

Tel: +1 (909) 280-6688; +86 138 1118 3009

Email: yqfu5166@ucla.edu | Website: <https://fyq5166.github.io/>

Education

University of California, Los Angeles (UCLA)

Master of Science in Computer Science

Los Angeles, CA

Sep. 2026 - Jun. 2028 (Expected)

University of Hong Kong (HKU)

BEng in Computer Science, minor in Finance

Departmental graduation rank: 10/133

Hong Kong, China

Sep. 2022 - Jun. 2026

University of California, Davis (UCD)

Exchange Student, Computer Science | GPA: 3.8/4.0

Davis, CA

Jan. 2025 - Jun. 2025

Publication

R. Cao, F. Lei, H. Wu, J. Chen, Y. Fu, et al. Spider2-V: How Far Are Multimodal Agents From Automating Data Science and Engineering Workflows? NeurIPS 2024, Datasets and Benchmarks Track. [Paper](#)

Research & Project Experience

ParaMind: Distributed LLM Inference Prototype [Code](#) · [dev branch](#)

Undergraduate Capstone Project

HKU

Jan. 2026 - May 2026

- Built a prototype to explore running a language model across resource-constrained machines, adapting Qwen and Llama loading, layer-wise execution, and shard-boundary inputs/outputs.
- Developed greedy layer placement and local replanning for node joins and departures, using measured matrix-multiplication performance and memory budgets to assign model layers to heterogeneous nodes.
- Integrated the inference runtime with an independently developed P2P subsystem through defined transport interfaces; demonstrated **two-node model loading and text generation** with both Qwen and Llama.

TABULA-R²: Table Reasoning Evaluation

Remote Research Assistant, supervised by Prof. Philipp Koehn, JHU

Remote

Aug. 2025 - Nov. 2025

- Investigated whether context-limited local LLMs could answer table questions through executable operations; constructed single-table, multi-table, and distractor tasks from 129 real-world OWID tables containing 70K+ rows.
- Defined a table-operation language with structured parsing and syntax validation; built a pandas-based executor returning results and errors to support multi-turn execution and correction.
- Compared models, prompts, and table inputs; observed weaker multi-table performance and **format-related execution failures** in this setup. Documented findings in a [technical report](#) and released [code and data](#).

WorldMap: Interactive Website Representation and Trajectory Generation

Research Assistant, supervised by Prof. Tao Yu

XLANG Lab, HKU

Jul. 2024 - Nov. 2024

- Built reusable website interaction graphs for agent grounding and planning, capturing components, action outcomes, and navigation; XPath-bound actions enabled direct interaction and replay of recorded trajectories.
- Generated **candidate computer-use training data** on WebArena through graph-guided LLM action selection, step-intent annotation, task summaries, and deduplication, without manually authoring each trajectory.

Spider2-V: Multimodal Data Workflow Benchmark [Paper](#)

Research Assistant, supervised by Prof. Tao Yu

XLANG Lab, HKU

Mar. 2024 - Jun. 2024

- Helped define enterprise data-workflow tasks to evaluate multimodal agents, assessing Airflow and Snowflake for practical relevance, GUI interaction, and feasibility of programmatic setup and evaluation.
- Built **executable tasks** for Airflow, LibreOffice Calc, and Google Search, including instructions, environment setup/reset, and automated success checks; ran complete Snowflake workflows and checked their evaluation scripts.
- Ran and monitored selected observation ablations and helped analyze results: Set-of-Mark (SoM) outperformed unaligned screenshot-plus-tree input on the evaluated subset.
- Prepared benchmark visualizations and assisted with manuscript drafting for the NeurIPS 2024 Datasets and Benchmarks paper.

NLP for Scam-Related Text Analysis

Independent Researcher, supervised by Kieren Lovell

Cambridge OSRP

Jun. 2023 - Sep. 2023

- Explored word-vector representations for scam-related text analysis, using NLTK tokenization, part-of-speech tagging, and lemmatization to prepare the text.

Internship

ByteDance · Volcano Engine, Ark Group

AI Agent Engineering Intern, Data Department

Beijing

Jun. 2025 - Aug. 2025

- Evaluation: Built a framework to test tool selection, execution, and final answers across local, API, and MCP tools, with configurable exact-match, regex, and LLM checks and structured logs; piloted internally.
- Cookbooks: Authored complementary Arkitect references to help developers and coding agents use framework interfaces: a tutorial notebook and a Python reference, both merged into the official repository through [PR #211](#).
- On-call: Built a workflow to reuse incident knowledge for repeated questions, retrieving from 1,000+ incident records and company knowledge via VikingDB, with Grafana dashboard links and TOS-stored chat images.
- Implemented **explicit human takeover and reply approval** so an LLM assistant could shift from answering to supporting responders; logged copy/send actions as adoption signals and piloted within a 30+ member team.

Technical Skills & Languages

Programming: Python (primary), C/C++, Java, R, SQL.

ML & Data: PyTorch, Transformers, pandas, NumPy; tool calling and LLM-assisted evaluation.

Systems: Linux, Docker, Git, FastAPI, SQLite, Electron.

Languages: Mandarin (Native), Cantonese (Intermediate), English (Fluent).