

Yeqiao Fu

Tel: +86 138 1118 3009; +852 5640 9355 | Email: yqfu5166@ucla.edu | Website: <https://fyq5166.github.io/>

EDUCATION

University of California, Los Angeles (UCLA)

Los Angeles, CA

Master of Science in Computer Science

Sept. 2026- Jun. 2028 (Expected)

University of Hong Kong (HKU)

Hong Kong, China

BEng in Computer Science, minor in Finance

Sept. 2022- Jun. 2026

Departmental graduation rank: 10/133

University of California, Davis (UCD)

Davis, CA

Exchange, major in Computer Science | GPA: 3.8/4.0

Jan. 2025- Jun. 2025

- **Relevant Coursework:** Artificial Intelligence, Data Structures and Algorithms, Probability and Statistics, Linear Algebra, Linux-Based Data Analysis, R-Based Data Visualization and Analysis

PUBLICATION

R. Cao, F. Lei, H. Wu, J. Chen, **Y. Fu**, et al. Spider2-V: How Far Are Multimodal Agents From Automating Data Science and Engineering Workflows? NeurIPS 2024, Datasets and Benchmarks Track.

Research Experience

Spider2-V: Multimodal Data Workflow Benchmark

XLANG Lab, HKU

Research Assistant, supervised by Prof. Tao Yu

Mar. 2024 - Jun. 2024

- Helped scope enterprise data-workflow tasks by exploring Airflow and Snowflake for practical relevance, GUI interaction, and programmatic setup and evaluation.
- Built executable benchmark tasks for Airflow, LibreOffice Calc, and Google Search, including instructions, environment setup/reset, and automated success checks; validated the Snowflake task set.
- Ran and monitored observation-space ablations, comparing screenshots, accessibility trees, and Set-of-Mark (SoM); helped analyze the mixed effects of richer inputs on the evaluated subset.
- Prepared benchmark visualizations and assisted with manuscript drafting for the NeurIPS 2024 Datasets and Benchmarks paper.

TABULA-R²: Table Reasoning Evaluation

Remote

Remote Research Assistant, supervised by Prof. Philipp Koehn, JHU

Nov. 2025 - Feb. 2026

- Curated 129 real-world OWID tables (70K+ rows) and designed single-table, multi-table, and distractor tasks through LLM-assisted generation with SQL, manual, and LLM-based answer checks.
- Built a constrained-DSL evaluation framework with PLAN/END parsing, syntax checks, and pandas execution, returning results and errors for multi-turn correction by local LLMs.
- Compared models, prompting, and table-observation settings; analyzed execution failures, documented findings in a technical report, and released code and data. [\[Code & Data\]](#)

ParaMind: Distributed LLM Inference Prototype

HKU

Undergraduate Capstone Project

Jan. 2026 - May 2026

- Implemented layer-sharded inference for Qwen and Llama, adapting model loading and execution to their architectures so each shard could run its assigned model components.
- Developed compute-aware, memory-constrained greedy layer placement and local replanning for node joins and departures, using measured matrix-multiplication performance to estimate node capability.
- Integrated the inference runtime with an independently developed P2P subsystem through defined transport interfaces; demonstrated model loading and text generation in a two-node prototype.

WorldMap: Interactive Website Representation and Trajectory Generation

XLANG Lab, HKU

Research Assistant, supervised by Prof. Tao Yu

Jul. 2024 - Nov. 2024

- Built reusable website interaction graphs capturing components, action outcomes, and navigation paths; XPath-bound actions supported direct component interaction and replay of recorded trajectories.
- Developed and ran a graph-guided pipeline on WebArena, using an LLM to generate trajectories, annotate steps, summarize tasks, and remove duplicates, producing candidate training data for computer-use agents.

NLP for Scam-Related Text Analysis

Cambridge OSRP

Independent Researcher, supervised by Kieren Lovell

Jun. 2023 - Sep. 2023

- Applied NLTK tokenization, part-of-speech tagging, and lemmatization to prepare scam-related text for exploratory analysis with word-vector representations.

INTERNSHIP

ByteDance · Volcano Engine, Ark Group

Beijing

AI Agent Engineering Intern, Data Department

Jul. 2025 - Aug. 2025

- Built a tool-use evaluation framework for local, API, and MCP tools, assessing calls, execution, and final answers with exact-match, regex, and LLM judges; piloted internally.
- Authored an Arkitect tutorial notebook for developers and a Python reference for coding agents; both merged into the official repository through [\[PR #211\]](#).
- Built an on-call retrieval workflow from 1,000+ incident records and company knowledge using VikingDB, with Grafana dashboard links and TOS-stored images for incident context.
- Implemented human takeover and reply approval for an LLM on-call assistant, logging copy/send actions as adoption signals; piloted within a 30+ member team.

Skillset & Interests

Languages: Mandarin (Native), Cantonese (Intermediate), English (Fluent)

Technical Skills: Python (primary), C/C++, Java, R, SQL; PyTorch, Transformers, pandas, NumPy; Linux, Docker, Git, FastAPI, SQLite, Electron; LLM / agent tooling (tool-calling, LLM-assisted evaluation); HTML/DOM and accessibility representations; structured trajectory logging; Matplotlib, Jupyter, LaTeX.